

```

# Importing necessary Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# Load the dataset
# Assuming 'house_prices.csv' is in the current directory
data = pd.read_csv('/content/house-prices MF ori.csv')
print("First 5 rows of the dataset:")
print(data.head())

# Dataset info
print("\nDataset Information:")
print(data.info())

# Check for missing values
print("\nMissing Values: ")
print(data.isnull().sum())

# Fill missing values (example: median for numerical columns)
data['SqFt'].fillna(data['SqFt'].median(), inplace=True)
data['Bedrooms'].fillna(data['Bedrooms'].median(), inplace=True)

# Handle outliers (example: capping)
upper_limit = data['Price'].quantile(0.95)
data['Price'] = np.where(data['Price'] > upper_limit, upper_limit, data['Price'])

# Encoding the 'Location' column using one-hot encoding
data = pd.get_dummies(data, columns=['Neighborhood', 'Home'], drop_first=True)

from sklearn.preprocessing import MinMaxScaler

# Normalize numerical columns
scaler=MinMaxScaler()
data[['SqFt', 'Bedrooms']] = scaler.fit_transform(data[['SqFt', 'Bedrooms']])

# Define features and target variable
X = data.drop('Price', axis=1)
y = data['Price']

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42) # Changed 'x' to 'X'
print(f"Training set size: {X_train.shape}")
print(f"Testing set size: {X_test.shape}")

# Initialize and train the model
model = LinearRegression()
model.fit(X_train, y_train)

# Display coefficients
print("Model Coefficients:", model.coef_)
print("Intercept: ", model.intercept_)

y_pred = model.predict(X_test)

# Evaluate the model
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
r2 = r2_score(y_test, y_pred)

print(f"RMSE: {rmse:.2f}")
print(f"R2: {r2:.2f}")

# Scatter plot of actual vs predicted prices
plt.figure(figsize=(8, 6))
plt.scatter(y_test, y_pred, alpha=0.7)
plt.title("Actual vs Predicted Prices")
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.show()

# Residual plot
residuals = y_test - y_pred
plt.figure(figsize=(8, 6))

```

```
sns.histplot(residuals, kde=True, bins=30, color='blue')  
plt.title("Distribution of Residuals")  
plt.xlabel("Residuals")  
plt.show()
```



First 5 rows of the dataset:

	Home	Price	SqFt	Bedrooms	Bathrooms	Offers	Brick	Neighborhood
0	1	114300	1790	2	2	2	2	101
1	2	114200	2030	4	2	3	2	101
2	3	114800	1740	3	2	1	2	101
3	4	94700	1980	3	2	3	2	101
4	5	119800	2130	3	3	3	2	101

Dataset Information:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 128 entries, 0 to 127

Data columns (total 8 columns):

#	Column	Non-Null Count	Dtype
0	Home	128 non-null	int64
1	Price	128 non-null	int64
2	SqFt	128 non-null	int64
3	Bedrooms	128 non-null	int64
4	Bathrooms	128 non-null	int64
5	Offers	128 non-null	int64
6	Brick	128 non-null	int64
7	Neighborhood	128 non-null	int64

dtypes: int64(8)

memory usage: 8.1 KB

None

Missing Values:

Home	0
Price	0
SqFt	0
Bedrooms	0
Bathrooms	0
Offers	0
Brick	0
Neighborhood	0

dtype: int64

Training set size: (102, 134)

Testing set size: (26, 134)

Model Coefficients: [ 3.96012842e+04 5.45124252e+03 8.13017724e+03 -6.45453118e+03

-1.27413706e+04 1.85749763e+04 -7.19045383e+03 -5.61674296e+03  
 -6.03471359e+03 -2.15627637e+04 -3.04023473e-12 -1.97446774e+04  
 1.10140751e+03 1.33781364e+03 -3.39578045e+03 -1.17084488e+04  
 0.00000000e+00 2.27373675e-12 2.41387912e+03 3.35975003e+03  
 1.72793723e+04 3.18371053e+03 9.12524628e+03 -2.74953356e+04  
 -3.63797881e-12 -3.63797881e-12 6.53745034e+03 1.04874619e+04  
 -7.29829897e+03 -1.01556054e+04 3.00434371e+03 1.95116618e+04  
 3.63797881e-12 9.09494702e-12 -2.49548010e+04 6.94529090e+03  
 1.25593798e+04 1.27329258e-11 -1.15135630e+04 -5.13533257e+03  
 9.85001160e+03 -3.57628397e+03 -9.09494702e-12 -6.61517308e+03  
 -7.41493051e+03 1.01027142e+03 5.45696821e-12 -1.27762136e+04  
 -9.19306254e+03 1.41710815e+04 5.72569978e+03 -9.32232069e-12  
 7.61268286e+03 -1.84037413e+03 -5.00021404e+03 -4.58704921e+03  
 1.79273095e+04 -1.19866728e+04 2.56367025e+03 1.75591339e+04  
 -1.63087802e+04 -3.63797881e-12 -8.52651283e-12 -1.51686991e+03  
 -1.29136640e+04 2.01494026e+03 1.59456673e+04 4.32881768e+01  
 1.72335828e+03 -6.09873348e+03 -9.09494702e-12 -3.39453858e+03  
 7.46231129e+03 2.36057001e+04 -1.45758537e+04 -1.81898940e-12  
 1.20328513e+04 1.78328246e+03 9.69121977e+03 -7.23175024e+03  
 -4.32599221e+03 -4.03072272e+03 -9.75537274e+03 8.55224079e+03  
 7.01699484e+03 9.83710532e+02 -1.81898940e-12 -3.18323146e-12  
 3.69725013e+03 1.91520138e+03 -1.81898940e-12 6.10484867e+03  
 1.04417796e+03 7.35544671e+03 5.77988014e+03 -4.04283690e+03  
 2.07492611e+03 -1.21227495e+04 -3.20820709e+03 2.27787310e+04  
 0.00000000e+00 -5.82371238e+03 0.00000000e+00 -1.39908751e+04  
 0.00000000e+00 5.36238417e+03 5.99461344e+03 1.88949021e+03  
 3.92426374e+03 5.95916581e+03 0.00000000e+00 0.00000000e+00  
 6.71526023e+02 -4.78677943e+03 -2.04816912e+04 -1.20786996e+04  
 8.14960135e+02 9.52735226e+03 -8.65206584e+03 0.00000000e+00  
 -9.20669202e+03 -1.19234496e+04 -1.16864593e+04 4.35743707e+03  
 0.00000000e+00 1.09835803e+04 -3.24850784e+03 -1.70560425e+04  
 1.92484674e+03 -1.46552740e+04 2.69577332e+03 0.00000000e+00  
 -7.36649834e+03 4.47279254e+03]

Intercept: 124620.53965201211

RMSE: 11240.03

R2: 0.78

<ipython-input-30-4aa918083afd>:25: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained a The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting



For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[co

```
data['SqFt'].fillna(data['SqFt'].median(), inplace=True)
```

<ipython-input-30-4aa918083afd>:26: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained a The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[co

```
data['Bedrooms'].fillna (data['Bedrooms'].median(), inplace=True)
```



Actual vs Predicted Prices